

Advanced Methodologies for Decoupled Risk Management in Multi-Agent Algorithmic Trading Systems

The Architectonic Crisis in Heterogeneous Multi-Agent Ensembles

The integration of heterogeneous alpha-generating agents into a unified algorithmic trading architecture presents profound challenges in risk management, capital allocation, and execution logic. Advanced ensembles increasingly rely on diverse, non-correlated agents to navigate shifting market regimes. A typical sophisticated triad might include a deep-learning foundation model predicting one-day forward returns via probabilistic quantiles, an Adaptive Kalman statistical filter isolating lagging price momentum, and a Deep Order Flow Imbalance (OFI) agent detecting real-time microstructural volume shocks. While combining predictive, lagging, and real-time models theoretically provides robust diversification, forcing these fundamentally distinct strategies to share execution parameters derived from a single model introduces a catastrophic structural flaw.

When an ensemble utilizes a deep learning predictive model—such as a Chronos AI agent—to forecast forward returns and generate strict confidence bounds (e.g., 0.1 and 0.9 quantiles), the resulting brackets represent a specific statistical expectation of volatility over a fixed duration. Applying these precise quantiles as universal Take-Profit (TP) and Stop-Loss (SL) parameters for fundamentally different models creates a severe temporal and logical mismatch. A Kalman-based momentum trade relies on the persistence of a trend and inherently requires wider physical boundaries to survive natural market retracements and cyclical pullbacks. Conversely, an OFI agent operates on microstructural timescales where the validity of an order book imbalance signal may expire within minutes, necessitating highly reactive, tight execution bounds¹.

Forcing conflicting strategies into the execution parameters of a single predictive model results in premature stop-outs for the lagging and microstructure agents. Their trades are liquidated not because their underlying alpha thesis was invalidated, but because the market breached a volatility quantile predicted by a completely different model. This phenomenon illustrates the "Monolithic Fallacy" in financial machine learning, wherein an architecture attempts to simultaneously optimize entry and exit logic within a single, rigidly coupled objective function, ignoring the asymmetric nature of market liquidity and volatility⁴.

In systems utilizing Bayesian Thompson Sampling to dynamically allocate capital based on historical win rates, these premature exits trigger a vicious cycle. Thompson Sampling algorithms update the posterior probability distributions of each agent based on observed rewards or penalties⁵. When the Kalman and OFI agents are systematically stopped out due to incompatible risk brackets, the bandit algorithm registers these events as failed predictions rather than execution failures. Consequently, the posterior distributions for these agents shift

negatively, causing the Thompson Sampling mechanism to mathematically starve them of capital. This dynamic reduces their portfolio allocation to nearly zero, collapsing a diversified multi-agent ensemble into a fragile, single-agent allocation⁵. Resolving this architectural deterioration requires completely decoupling the alpha generation layer from the risk management and execution layer.

Institutional Decoupling of Alpha and Execution

Top-tier institutional quantitative funds manage the integration of fundamentally conflicting models by strictly separating signal generation from trade liquidation. This decoupling ensures that the evaluation of a predictive model's accuracy is not polluted by the mechanical frictions of order execution.

The Hunter-Guard Paradigm

The "Hunter-Guard" architecture serves as a premier example of this necessary decoupling, finding a compelling analog in the Dual Process Theory of human cognition⁴. In this framework, alpha agents act exclusively as "Hunters" (System 2: slow, deliberative, analytical). Their sole objective is to identify divergences from equilibrium where the probability of price appreciation exceeds transaction costs, generating a directional signal or target portfolio weight⁴. They are optimized for precision and are largely unconcerned with the mechanics of the exit, focusing entirely on capturing idiosyncratic alpha.

The execution and risk management layer acts as the "Guard" (System 1: fast, reactive). The Guard model operates independently of the entry logic, evaluating the portfolio purely based on prevailing market conditions, downside deviation, systemic beta factors, and microstructural deterioration⁴. By separating the initiation logic from the liquidation logic, the system allows for the optimization of asymmetric risk-reward profiles. The Hunter can be tuned to be patient, waiting for high-conviction setups, while the Guard can be tuned to be highly reactive, prioritizing the capture of necessary exit signals to prevent catastrophic loss⁴. This asymmetry—constrained entry combined with unconstrained, agnostic exit logic—creates a "ratchet" effect that is crucial for institutional capital preservation⁴.

Continuous Adaptive Pipelines

Modern multi-agent frameworks further emphasize this separation through continuous, adaptive pipelines. Architectures like AlphaCrafter utilize a multi-agent system where factor discovery is continuously calibrated against prevailing market regimes before being passed to an execution agent⁸. The execution agent translates the aggregated signal ensemble into actual orders subject to explicit portfolio construction and risk constraints, dynamically reweighting without requiring the alpha generators to understand the execution friction⁹. Similarly, the TiMi (Trade in Minutes) system architecturally decouples strategy development from minute-level deployment. This paradigm separates complex reasoning from time-sensitive execution, enabling comprehensive strategy development offline while maintaining quantitative-level efficiency and low latency across active market dynamics during deployment¹⁰.

To achieve an agent-agnostic risk management layer within a Thompson Sampling ensemble, the execution parameters must be derived directly from the underlying asset's current volatility regime and the immediate liquidity profile of the limit order book, rather than from the subjective prediction of a single constituent agent. The execution layer must transition from arbitrary geometric thresholds or coupled quantiles to adaptive, data-driven boundaries.

Macro-Variance Forecasting for Agent-Agnostic Boundaries

Calculating unbiased, dynamic TP/SL bounds requires synthesizing macroeconomic variance expectations with microstructural realities. The foundational layer for any dynamic exit calculation is a robust estimate of the asset's variance over the intended holding period. Two primary methodologies dominate this space: Average True Range (ATR) parameterization and Generalized Autoregressive Conditional Heteroskedasticity (GARCH) forecasting.

Average True Range (ATR) Dynamics

The Average True Range (ATR) is a ubiquitous, asset-agnostic volatility indicator that measures the average absolute price movement range over a specified lookback period¹¹. Developed by J. Welles Wilder, the true range accounts for gaps between trading sessions by selecting the greatest of three values: the current high minus the current low, the absolute value of the current high minus the previous close, or the absolute value of the current low minus the previous close¹².

ATR-based trailing stops satisfy a monotonic non-decreasing ratcheting property, providing a structural, closed-form per-trade loss bound that naturally adapts to shifting volatility¹⁴. When implemented dynamically, an ATR multiplier expands the survival bracket during periods of market turbulence and tightens it during consolidation, helping to avoid stop-outs caused by

localized market noise¹³. For example, a system might utilize a $1.5 \times ATR$ distance for a stop-loss and a $3.0 \times ATR$ distance for a take-profit, inherently targeting a positive risk-reward asymmetry¹³.

Extensive empirical evaluations of ATR parameterization in autonomous trading agent swarms demonstrate that exit design meaningfully impacts risk-adjusted performance. Studies analyzing historical trades across vast hyperparameter grids indicate that dynamic ATR-based exits, combined with circuit-breaker overlays, can yield Sharpe ratio improvements of over 56% relative to static, heuristic-based baseline configurations¹⁷. Furthermore, analyses of point-in-time equity universes reveal a "stop-multiplier-insensitive plateau," where a range of ATR configurations produces consistently robust Sharpe ratios and uniform maximum consecutive negative-year runs, establishing the ATR multiplier as a structurally resilient metric for bounding risk¹⁴.

Despite its computational efficiency and effectiveness for trend-following strategies, ATR possesses a critical limitation for advanced quantitative systems: it is a purely symmetric and lagging measure. It assumes that upward volatility and downward volatility behave identically

and relies entirely on historical price action without forecasting future variance distributions¹³.

Econometric Variance Forecasting: The GARCH Family

Financial equities exhibit a pronounced "leverage effect," wherein negative price shocks increase subsequent volatility significantly more than positive shocks of the same magnitude²⁰. Standard historical measures like ATR, and even foundational Generalized Autoregressive Conditional Heteroskedasticity (GARCH(1,1)) models, treat positive and negative shocks identically, rendering them suboptimal for precise downside risk management¹⁹. GARCH(1,1) models variance as a recursive linear function of lagged conditional variances and lagged squared returns, capturing volatility clustering and mean reversion, but failing to capture directionality²³.

To capture this asymmetry, institutional risk models frequently employ the GJR-GARCH (Glosten-Jagannathan-Runkle) extension or Threshold GARCH (TGARCH) models. The GJR-GARCH model introduces an additional parameter to account specifically for the

asymmetric response to negative shocks²¹. The conditional variance σ_t^2 is modeled mathematically as:

$$\sigma_t^2 = \omega + \alpha \epsilon_{t-1}^2 + \gamma I_{[\epsilon_{t-1} < 0]} \epsilon_{t-1}^2 + \beta \sigma_{t-1}^2$$

In this formulation, ω represents the long-run baseline variance constant. The parameter α represents the symmetric reaction to the previous period's shock, known as the ARCH term.

The parameter β represents the persistence of volatility, or the GARCH term. The critical innovation is γ , which captures the asymmetric leverage effect. The indicator function

$I_{[\epsilon_{t-1} < 0]}$ equals 1 strictly if the previous shock was negative, thereby adding the γ coefficient to the volatility calculation and amplifying the forecasted variance in response to downside moves²¹.

Alternative formulations, such as logarithmic GARCH (logGARCH), apply a logarithmic transformation to volatility shocks to improve numerical stability, interpretability, and tail control, ensuring that the conditional variance remains strictly positive without requiring artificial parameter constraints²⁰. Threshold GARCH models extend this further by allowing the parameters to change completely based on whether the market is in a positive or negative regime, effectively capturing the leptokurtic, persistent, and asymmetric features inherent in financial time series²⁰.

For an agent-agnostic risk manager, replacing predictive deep learning quantiles with a GJR-GARCH forecasted volatility band provides a mathematically rigorous, unbiased baseline for the asset's expected variance over the holding period²⁸. This econometric approach isolates the asset's innate statistical behavior from the alpha generation logic, ensuring that a Kalman filter's momentum trade is given exactly the mathematical breathing room the asset's current

volatility profile dictates.

Feature Profile	Average True Range (ATR)	GJR-GARCH / TGARCH Forecasting
Fundamental Calculation	Rolling simple or exponential average of absolute price extremes.	Econometric modeling of conditional variance based on past innovations.
Directional Symmetry	Strictly symmetric; treats upward and downward ranges equally.	Asymmetric; explicitly accounts for the equity leverage effect via indicator functions.
Temporal Perspective	Lagging indicator based on a predefined historical lookback window.	Predictive conditional variance forecast for the $t + 1$ horizon.
Computational Overhead	Extremely low ($O(1)$ updates per time step).	Moderate to High; requires Maximum Likelihood Estimation (MLE) optimization.
Optimal System Application	Adaptive trailing stops and immediate intraday volatility scaling.	Baseline daily risk budgeting, regime identification, and core boundary setting.

Microstructural Liquidity Profiling and Execution

While GJR-GARCH provides an optimal macroeconomic variance band, placing a stop-loss precisely at a calculated standard deviation boundary (e.g., $P_{entry} - 1.5\sigma$) creates severe execution vulnerability. Financial markets are driven by limit order books (LOB), and setting stops at arbitrary statistical distances entirely ignores the physical location of liquidity.

The Danger of Price Cascades and Stop Runs

Stop-loss orders generate positive-feedback trading, often accelerating trends and contributing to rapid, self-reinforcing price movements known as price cascades or stop runs³⁰. When a statistical boundary happens to fall in a "liquidity void"—an area of the LOB with

minimal resting limit orders—a triggered market-sell order will slice through multiple price levels, resulting in severe slippage³. Institutional market makers actively hunt these statistical thresholds, pushing prices into liquidity voids to trigger retail and algorithmic stop-losses before absorbing the resulting aggressive flow³².

The empirical evidence demonstrates that currency and equity stop-loss orders contribute significantly to these price discontinuities. The response to stop-loss orders is structurally larger and lasts longer than the response to take-profit orders, which generate negative-feedback trading and typically absorb volatility³⁰. Therefore, a mathematically sound risk boundary must be contextually adjusted to survive the microstructure of the order book.

Order Flow Imbalance (OFI) and Order Book Depth

To translate a theoretical volatility forecast into a safe, executable bracket order, the execution layer must continuously analyze Order Book Imbalance (OBI) and microstructural depth¹. OBI measures the disparity between resting buy (bid) and sell (ask) limit orders at varying price levels, serving as a bridge between the static picture of the book and the dynamic process of price formation³.

The Static Order Book Imbalance (SOBI) provides a normalized metric of this disparity and is defined mathematically as:

$$SOBI = \frac{\sum_{i=1}^N Q_{bid,i} - \sum_{i=1}^N Q_{ask,i}}{\sum_{i=1}^N Q_{bid,i} + \sum_{i=1}^N Q_{ask,i}}$$

In this equation, Q represents the resting quantity at depth level i , aggregated up to a specified depth level N ³⁶. A positive SOBI indicates strong bid-side dominance and support, while a negative SOBI indicates ask-side dominance and heavy resistance³. Relying on the traditional top-of-book mid-price to place stops is insufficient in imbalanced markets. Instead, quantitative systems calculate the Volume Adjusted Mid Price (VAMP) or the effective execution price to determine the true center of market mass³⁶. The VAMP crosses price and quantity between the bid and ask sides:

$$VAMP_{bbo} = \frac{P_{bid} \times Q_{ask} + P_{ask} \times Q_{bid}}{Q_{bid} + Q_{ask}}$$

By extending this logic to N levels of depth, algorithms gain insight into short-term supply and demand imbalances, identifying where the bulk of resting liquidity actually resides¹.

The Hybrid Execution Formulation: Tucking Behind Liquidity

A state-of-the-art risk execution agent must combine econometric forecasting with microstructural mapping to establish "absorption walls" or liquidity nodes. Instead of placing the stop-loss blindly at the exact GJR-GARCH threshold, the system must scan the LOB depth

profile in the immediate vicinity of that threshold¹.

If the statistical threshold falls in a thin liquidity zone, the execution agent moves the physical stop-loss slightly wider, "tucking" it securely behind a massive resting bid wall. Because market makers and institutional participants defend these high-liquidity zones, they act as microstructural support, absorbing transient aggressive flow³¹. If this wall is breached, the original trade thesis is definitively invalidated, justifying the exit.

This hybrid process operates sequentially:

1. **Macro-Variance Calculation:** Calculate the $t + 1$ expected conditional variance σ_{t+1}^2 using a rolling GJR-GARCH or TGARCH model to account for asymmetric leverage effects²².
2. **Base Bracket Generation:** Establish the theoretical Stop-Loss (SL_{base}) and Take-Profit (TP_{base}) using a multiple of the forecasted standard deviation, adjusted for the asset's current drift and the specific alpha agent's expected holding period.
3. **Microstructural Adjustment:** Scan the Limit Order Book depth within a predefined fractional radius (ϵ) of SL_{base} and TP_{base} .
4. **Liquidity Anchoring:** Identify the local maximum of cumulative resting liquidity (L_{max}) within the radius. Adjust the final executable brackets to rest immediately behind these structural liquidity nodes, protecting the orders from transient stop-hunting and liquidity voids³.

This decoupled execution agent serves the Chronos, Kalman, and OFI models equally. It allows momentum and volume strategies to survive natural asset volatility while exiting strictly when empirical market structure breaks down, entirely independent of the Chronos model's specific predictive quantiles.

Rectifying Bayesian Thompson Sampling Capital Starvation

Decoupling the execution layer solves the mechanical issue of premature stop-outs, but the multi-agent ensemble must also address the mathematical deterioration of the Thompson Sampling allocation. If the bandit algorithm continues to evaluate the alpha agents based on a flawed, coupled reward function (e.g., the realized PnL of trades that were previously mismanaged by misaligned brackets), capital allocation will remain distorted.

The Reward Function Misalignment

In standard Thompson Sampling for Multi-Armed Bandits (MAB), an agent selects an action that maximizes the expected reward with respect to a randomly drawn belief from a posterior distribution⁶. For binary outcomes (win/loss), this is modeled using Beta distributions

$\text{Beta}(\alpha, \beta)$, updated via simple success/failure observations³⁸. For continuous rewards,

Gaussian priors updated via maximum likelihood estimation are utilized⁴¹.

The architectural flaw in using realized trade PnL to update the bandit is that it inextricably entangles the alpha generation capability with the execution friction. If the OFI agent correctly identifies a short-term liquidity shock, but the universal execution layer subsequently mismanages the exit due to an arbitrary hold time or a poor spread crossing, the OFI agent is penalized with a "loss." The Thompson Sampling algorithm updates the OFI agent's posterior distribution negatively, lowering its expected value and ultimately starving it of capital⁵. To restore parity, the ensemble must redefine its reward function to measure the pure predictive power of the underlying alpha agents, entirely separate from the execution layer's PnL.

AlphaEval Metrics: Information Coefficient (IC)

Quantitative evaluation frameworks accomplish this by focusing on continuous, execution-agnostic metrics like the Information Coefficient (IC) and Rank Information Coefficient (RankIC)⁴³.

The Information Coefficient is defined as the Pearson correlation between the agent's generated alpha score (or target weight) and the actual forward realized return of the asset across the defined holding period⁴³. The mathematical representation across N assets is:

$$IC_t = \frac{\sum_{i=1}^N (S_{t,i} - \bar{S}_{t,:})(y_{t,i} - \bar{y}_{t,:})}{\sqrt{\sum_{i=1}^N (S_{t,i} - \bar{S}_{t,:})^2} \sqrt{\sum_{i=1}^N (y_{t,i} - \bar{y}_{t,:})^2}}$$

Where $S_{t,i}$ represents the signal score for asset i at time t , and $y_{t,i}$ represents the realized return. RankIC utilizes the Spearman rank correlation, evaluating how well the agent ranks a cross-section of assets, providing critical stability against outliers⁴³.

By updating the Thompson Sampling posterior distributions using a composite Predictive Power Score (PPS) derived from IC and RankIC instead of binary trade survival, the ensemble explicitly rewards the Chronos, Kalman, and OFI agents for their mathematical accuracy⁴³. The Kalman agent is rewarded when its momentum thesis proves statistically valid over its intended time horizon, regardless of whether the physical bracket order experienced microstructural slippage. This strictly aligns the bandit's exploration-exploitation mechanism with pure alpha generation.

Risk-Aware Bandits: Mean-Variance Thompson Sampling (MVTs)

Standard Thompson Sampling seeks solely to maximize the expected reward, entirely ignoring the variance (risk) of the distribution, which is a critical oversight in financial applications⁴⁰.

Mean-Variance Thompson Sampling (MVTs) directly integrates risk tolerance into the exploration-exploitation algorithm, ensuring that agents generating highly volatile, erratic signals are penalized in the allocation weighting⁴⁰.

MVTs models the reward using the Markowitz mean-variance objective:

$$MV_i = \rho\mu_i - \sigma_i^2$$

Where μ_i is the expected return (or IC) of agent i , σ_i^2 is the variance of agent i 's returns, and ρ is the system's defined risk tolerance parameter⁴⁰. The algorithm requires sampling both a mean parameter and a precision parameter to capture the risk-adjusted value of each arm. At each time step $t \geq K + 1$, MVTs samples a precision parameter $\tau_{i,t}$ from a Gamma distribution:

$$\tau_{i,t} \sim \text{Gamma}(\alpha_{i,t-1}, \beta_{i,t-1})$$

It then samples a mean parameter $\theta_{i,t}$ using the updated history from a Gaussian distribution:

$$\theta_{i,t} \sim \mathcal{N}\left(\hat{\mu}_{i,t-1}, \frac{1}{T_{i,t-1}}\right)$$

The ensemble then selects the agent that maximizes the sampled mean-variance combination:

$$i(t) = \arg \max_{i \in [K]} \left(\rho\theta_{i,t} - \frac{1}{\tau_{i,t}} \right)$$

This provides a mathematically sound mechanism for balancing the conflicting outputs of the Chronos, Kalman, and OFI models. If the Chronos agent consistently produces high returns but with massive variance, the σ_i^2 penalty will reduce its overall allocation weight, allowing more consistent agents to receive capital⁴⁰.

Thompson Sampling Variant	Objective Function	Key Mechanism	Optimal Use Case
Standard Bernoulli TS	Maximize cumulative expected successes.	Beta posterior updates based on binary win/loss metrics.	Basic A/B testing; simple classification tasks.
Gaussian TS	Maximize cumulative expected continuous returns.	Gaussian posterior updates based on continuous reward metrics.	Environments with normally distributed, stable rewards.

Mean-Variance TS (MVTs)	Maximize risk-adjusted return ($MV_i = \rho\mu_i - \sigma_i^2$).	Samples both precision (Gamma) and mean (Gaussian) to penalize high-variance arms.	Financial portfolio allocation; risk-averse algorithmic trading.
Combinatorial Adaptive Discounted TS (CADTS)	Maximize recent risk-adjusted return under non-stationarity.	Applies sliding windows and geometric discounting to decay obsolete historical data.	Shifting market regimes (Bull/Bear/Chop transitions).

Addressing Non-Stationarity with CADTS

Because the efficacy of the three agents varies violently depending on the market regime (Bull, Bear, Chop), historic win rates rapidly become obsolete⁴⁶. Standard MVTs treats all historical data points with equal weight. Combinatorial Adaptive Discounted Thompson Sampling (CADTS) addresses concept drift by applying a sliding window and a discount factor to past observations⁴⁶.

By geometrically discounting older IC scores and focusing the posterior updates heavily on recent correlation metrics, CADTS allows the ensemble to rapidly pivot capital away from the lagging Kalman filter during choppy mean-reversion regimes, and heavily weight the OFI and Chronos agents. When a macro trend establishes, the discount factor quickly restores capital to the Kalman agent as its recent IC metrics surge⁴⁶.

Conclusion

The resolution to the multi-agent execution conflict lies in absolute architectural decoupling and statistical refinement. When predictive deep-learning models, lagging statistical filters, and real-time microstructure agents share a unified execution layer, mapping the predictive quantiles of one agent onto the others guarantees operational failure and capital starvation. To achieve robust performance, the architecture must transition to an agent-agnostic "Hunter-Guard" paradigm. The execution layer must calculate objective volatility bounds using GJR-GARCH to account for asymmetric equity leverage effects, and cross-reference these boundaries with real-time Limit Order Book depth. Tucking physical Stop-Loss and Take-Profit orders behind microstructural liquidity walls prevents slippage and stop cascades. Simultaneously, the Thompson Sampling ensemble must cease evaluating agents based on realized trade PnL. By scoring agents purely on their Information Coefficient (IC) relative to forward returns, and upgrading the allocation algorithm to Mean-Variance Thompson Sampling (MVTs) with adaptive discounting (CADTS), the system explicitly rewards pure predictive power

and penalizes erratic variance. This comprehensive decoupling ensures that every alpha generation model is granted a mathematically sound, volatility-normalized opportunity to realize its statistical edge.

Works cited

1. Order Book Features for Trading Insights | PDF | Applied Mathematics | Computing - Scribd, <https://www.scribd.com/document/945310331/Order-Book-Features-Extracted-by-Quantitative-Strategies>
2. Stochastic Price Dynamics in Response to Order Flow Imbalance: Evidence from CSI 300 Index Futures - IDEAS/RePEc, <https://ideas.repec.org/p/arx/papers/2505.17388.html>
3. What Is Order Book Imbalance? | Cube Exchange, <https://www.cube.exchange/what-is/order-book-imbalance>
4. (PDF) Decoupling Alpha and Risk: A Comprehensive Analysis of the Hunter-Guard Architecture and Stochastic Regularization in Financial Machine Learning - ResearchGate, https://www.researchgate.net/publication/398234740_Decoupling_Alpha_and_Risk_A_Comprehensive_Analysis_of_the_Hunter-Guard_Architecture_and_Stochastic-Regularization_in_Financial_Machine_Learning
5. Algorithm for Adaptive Experiments that Trade-off Statistical Analysis with Reward: Adaptively Combining Uniform Random Assignment and Thompson Sampling - arXiv, <https://arxiv.org/html/2112.08507>
6. A Tutorial on Thompson Sampling - Stanford University, https://web.stanford.edu/~bvr/pubs/TS_Tutorial.pdf
7. Multi-Armed Bandit (MAB) Methods in Trading - DayTrading.com, <https://www.daytrading.com/multi-armed-bandit>
8. AlphaCrafter: A Full-Stack Multi-Agent Framework for Cross-Sectional Quantitative Trading, https://www.researchgate.net/publication/404626683_AlphaCrafter_A_Full-Stack_Multi-Agent_Framework_for_Cross-Sectional_Quantitative_Trading
9. AlphaCrafter: A Full-Stack Multi-Agent Framework for Cross-Sectional Quantitative Trading, <https://arxiv.org/html/2605.05580v1>
10. Trade in Minutes! Rationality-Driven Agentic System for Quantitative Financial Trading - arXiv, <https://arxiv.org/html/2510.04787v1>
11. Why does my AI keep suggesting me to use ATR as an indicator for my stops? - Reddit, https://www.reddit.com/r/algotrading/comments/1mqcbu9/why_does_my_ai_keep_suggesting_me_to_use_atr_as/
12. ATR Indicator (Average True Range): Formula, Strategy & How to Use in 2026 | LiteFinance, <https://www.litefinance.org/blog/for-beginners/best-technical-indicators/atr-indicator/>
13. ATR: Technical Indicator | FTMO Academy,

- <https://academy.ftmo.com/lesson/atr-technical-indicator/>
14. Point-in-Time Backtesting of Momentum-Trend Equity Strategies: A Formal Bias Taxonomy, ATR Trailing Stop Analysis, and Investor - Preprints.org, https://www.preprints.org/frontend/manuscript/b8db286b07452ae7019c308413c28e10/download_pub
 15. Dynamic ATR-Based Stop-Loss and Take-Profit EMA Crossover Trend Quantification Strategy | by FMZQuant | Medium, <https://medium.com/@FMZQuant/dynamic-atr-based-stop-loss-and-take-profit-ema-crossover-trend-quantification-strategy-da87863a57c4>
 16. Average True Range: Dynamic Stop Loss Levels - LuxAlgo, <https://www.luxalgo.com/blog/average-true-range-dynamic-stop-loss-levels/>
 17. [2604.27150] Optimal Stop-Loss and Take-Profit Parameterization for Autonomous Trading Agent Swarm - arXiv, <https://arxiv.org/abs/2604.27150>
 18. Optimal Stop-Loss and Take-Profit Parameterization for Autonomous Trading Agent Swarm, <https://arxiv.org/html/2604.27150v1>
 19. Volatility Prediction in Trading: Models, Indicators, and Machine Learning, <https://arongroups.co/technical-analyze/volatility-prediction-in-trading/>
 20. The $\log TG - \$SV$ model: A threshold-based volatility framework with logarithmic shocks for exchange rate dynamics - AIMS Press, <https://www.aimspress.com/aimspress-data/math/2025/8/PDF/math-10-08-870.pdf>
 21. Analysis of Asymmetric GARCH Volatility Models with Applications to Margin Measurement - à www.publications.gc.ca, https://publications.gc.ca/collections/collection_2018/banque-bank-canada/FB3-5-2018-21-eng.pdf
 22. GJR-GARCH Volatility Documentation - V-Lab, <https://vlab.stern.nyu.edu/docs/volatility/GJR-GARCH>
 23. Machine Learning for Algorithmic Trading — Part 9 Time Series Models for Volatility Forecasting & Statistical Arbitrage | by Connie Zhou | Medium, <https://medium.com/@conniezhou678/machine-learning-for-algorithmic-trading-part-9-time-series-models-for-volatility-forecasting-c8f1afb4abd5>
 24. A Comparative Study of GARCH (1,1) and Black-Scholes Option Prices - McGill University, <https://www.mcgill.ca/desautels/files/desautels/Gar96oct.pdf>
 25. Volatility Forecasting: GARCH(1,1) Model - Portfolio Optimizer, <https://portfoliooptimizer.io/blog/volatility-forecasting-garch11-model/>
 26. The Threshold GARCH Model: Estimation and Density Forecasting for Financial Returns* - SciSpace, <https://scispace.com/pdf/the-threshold-garch-model-estimation-and-density-forecasting-ng4dn4va7y.pdf>
 27. A Threshold GARCH Model for Chilean Economic Uncertainty - MDPI, <https://www.mdpi.com/1911-8074/16/1/20>
 28. Forecasting Stock Market Volatility: A Wavelet-Enhanced Hybrid GARCH-Deep Learning Approach, https://jemsc.gom.ac.ir/article_4295_716a631a2a62194057eded67f3a86508.pdf
 29. DAY TRADING THE EURO-DOLLAR WITH A NEWS-BASED MODEL OF EXCHANGE

- RATES (*) - GRETA.it, https://www.greta.it/old/jae/poster/14_1_Tiveгна.pdf
30. Stop-Loss Orders and Price Cascades in Currency Markets | Request PDF - ResearchGate, https://www.researchgate.net/publication/314384055_Stop-Loss_Orders_and_Price_Cascades_in_Currency_Markets
 31. Inside The Market: Order Books And What You're Missing Out On - Bookmap, <https://bookmap.com/blog/inside-the-market-order-books-and-what-youre-missing-out-on>
 32. Liquidity Structure & Order Flow | Trading Indicator - LuxAlgo, <https://www.luxalgo.com/library/indicator/liquidity-structure-order-flow/>
 33. Order Flow Analysis: Tools, Strategies & Live Signals - StockGro, <https://www.stockgro.club/stoxo/resources/order-flow-analysis/>
 34. Order Book Imbalance | QuestDB, <https://questdb.com/glossary/order-book-imbalance/>
 35. How Order Flow Imbalance Can Boost Your Trading Success - Bookmap, <https://bookmap.com/blog/how-order-flow-imbalance-can-boost-your-trading-success>
 36. Market Making with Alpha - Order Book Imbalance - HftBacktest, <https://hftbacktest.readthedocs.io/en/latest/tutorials/Market%20Making%20with%20Alpha%20-%20Order%20Book%20Imbalance.html>
 37. Thompson sampling - Wikipedia, https://en.wikipedia.org/wiki/Thompson_sampling
 38. Thompson sampling - Andy Jones, <https://andrewcharlesjones.github.io/journal/thompson-sampling.html>
 39. Thompson Sampling for Bandits: A Comprehensive Guide for 2025 - Shadecoder, <https://www.shadecoder.com/topics/thompson-sampling-for-bandits-a-comprehensive-guide-for-2025>
 40. Thompson Sampling Algorithms for Mean-Variance Bandits - Proceedings of Machine Learning Research, <http://proceedings.mlr.press/v119/zhu20d/zhu20d.pdf>
 41. ICML Poster Connecting Thompson Sampling and UCB: Towards More Efficient Trade-offs Between Privacy and Regret, <https://icml.cc/virtual/2025/poster/45953>
 42. Adaptive Variance Inflation in Thompson Sampling: Efficiency, Safety, Robustness, and Beyond - OpenReview, <https://openreview.net/pdf?id=IBhxlNfxv>
 43. AlphaEval: A Comprehensive and Efficient Evaluation Framework for Formula Alpha Mining - arXiv, <https://arxiv.org/html/2508.13174v2>
 44. A Survey of Risk-Aware Multi-Armed Bandits - IJCAI, <https://www.ijcai.org/proceedings/2022/0784.pdf>
 45. Risk-averse Contextual Multi-armed Bandit Problem with Linear Payoffs - arXiv, <https://arxiv.org/pdf/2206.12463>
 46. Improving Portfolio Optimization Results with Bandit Networks - arXiv, <https://arxiv.org/html/2410.04217v2>